



[VNAppMob](#) / [Blogs](#)

/ Hosting multiple sites or applications using Docker and NGINX reverse proxy with Letsencrypt SSL

Hosting multiple sites or applications using Docker and NGINX reverse proxy with Letsencrypt SSL

So I decided to rehost my homepage and a couple other web pages and apps on a new server. And since we've been using Docker and Compose for some projects at work, I thought, hey, this could be a neat clean setup for multiple apps hosted on a single machine without installing a ton of local

So I decided to rehost my homepage and a couple other web pages and apps on a new server. And since we've been using Docker and Compose for some projects at work, I thought, hey, this could be a neat clean setup for multiple apps hosted on a single machine without installing a ton of local dependencies and managing everything by hand. Here's the story of how this went down, what I learned along the way, and how you can build the same setup without doing the same mistakes that I initially did.

Getting the hardware

First thing you need to run anything on the web is a server of some sort. Yes, yes, serverless is all the rage right now, but we're not going there.

There are all sorts of server setups available for all sorts of purposes and with all sorts of price tags. As a first step, I compiled a basic list of requirements for the kind of hosting I thought I'd need. It would be serving mostly static pages, with a few more dynamic apps with some background processes doing data collection and talking to some APIs, but there'd be no significant processing or other heavy lifting in there. So these were the basic min specs for the server I came up with:



- 20GB or more storage (*I'll need to serve some static content, so less than that won't do*)
- IPv6 enabled (*necessary for some of my projects that I might tell about at some point*)
- recent Linux kernel that'll support Docker properly (*surprisingly common issue with cheap VPS setups, as I found out*)

After researching the market a bit, I arrived at the conclusion that an entry-level VPS at a local provider here in Germany is probably what I want. There are many providers available out there and they change up their offerings rather fast, so I won't name any specific ones here - hit me up on Twitter if you want to know which one I ended up with.

A noteworthy alternative to a VPS is Amazon EC2. I actually used their free-tier `t2.micro` instance for a while to test an early iteration of this very setup, and it did the job perfectly. Of course, those are only free for 12 months - but if you just want to test the setup and possibly move to a different server later (or continue running on EC2, although that's considerably more expensive), it's something I can recommend.

DNS

You can run the applications under any domains or subdomains you like, provided their DNS is pointed to the server you have set up.

I personally run all applications as subdomains (and root domain) of `olex.biz`. The DNS is setup with A records for `olex.biz` and `*.olex.biz` pointing directly to the server's IP. This way I can choose any subdomains for new apps I want to deploy, and don't have to update DNS records every single time.

Setting up the software


Base OS and Docker

For my VPS, I chose to run an Ubuntu 18.04 image offered by the VPS provider. You may choose whichever distro you're familiar with - the only requirement is that it can run Docker. Get it installed and configured as you would configure any Internet-facing server - SSH with public key auth only, no root login, fail2ban, the usual setup. Important thing is at the end you have a working server with Docker and Compose available.

NGINX reverse proxy and SSL

The NGINX reverse proxy is the key to this whole setup. Its job is to listen on external ports 80 and 443 and connect requests to corresponding Docker containers, without exposing their inner workings or ports directly to the outside world. Additionally, with the SSL companion container the



 Diagram of using NGINX proxy on the server to talk to a dockerized application

For this job, I use [jwilder/nginx-proxy](#) and [docker-letsencrypt-nginx-proxy-companion](#).

My `docker-compose.yml` for them is placed in `/home/olex/proxy` and looks like this:

```
version: '2'

services:
  nginx-proxy:
    image: jwilder/nginx-proxy
    container_name: nginx-proxy
    ports:
      - "80:80"
      - "443:443"
    volumes:
      - conf:/etc/nginx/conf.d
      - vhost:/etc/nginx/vhost.d
      - dhparam:/etc/nginx/dhparam
      - certs:/etc/nginx/certs:ro
      - /var/run/docker.sock:/tmp/docker.sock:ro
    networks:
      - proxy
    restart: always

  letsencrypt:
    image: jrcs/letsencrypt-nginx-proxy-companion
    container_name: nginx-proxy-le
    volumes_from:
      - nginx-proxy
    volumes:
      - certs:/etc/nginx/certs:rw
      - /var/run/docker.sock:/var/run/docker.sock:ro
    restart: always

volumes:
  conf:
  vhost:
  dhparam:
  certs:
```



EXAMPLE.

name: nginx-proxy

Let's go through some details here to understand what's going on. There's two services, `nginx-proxy` and `letsencrypt`, as well as some volumes and a custom external network specified here.

nginx-proxy

`nginx-proxy` has a couple things happening:

- Its ports 80 and 443 are forwarded to the host, making it Internet-facing. No other containers we run on this machine will need their ports forwarded, all communication from and to the outside will be proxied through here - hence "reverse proxy".
- Various NGINX configuration directories are mounted as named volumes to keep them persistent on the host system. Those volumes are defined further down in the file.
- `/var/run/docker.sock` from the host is mounted. This allows the proxy to listen in to other containers starting and stopping on the host, and configure NGINX forwarding as needed. Containers need to present their desired hostnames and ports as environment variables that the proxy can read - more on that further below.
- Finally, the container is assigned to a `proxy` external network, which is described below.

letsencrypt

`letsencrypt` is a companion container to `nginx-proxy` that handles all the necessary SSL tasks - obtaining the required certificates from Let's Encrypt and keeping them up-to-date, and auto-configuring `nginx-proxy` to transparently encrypt all proxied traffic to and from application containers.

It's connected to `nginx-proxy` by sharing its volumes (`volumes-from:` directive). It also listens in on the host's `/var/run/docker.sock` to be notified when application containers are started, and to get the information from them to obtain the necessary SSL certificates.

Note that this container doesn't need to be put in the external network - it gets by only using the shared volumes and Docker socket, and never needs to talk to the outside world or another container directly.

External network

Last but not least, there's the external network. To understand why you might need it, you need to know how docker-compose handles networks by default:



can talk to each other and to the Internet.

- We want to deploy multiple applications on this server using Compose, each with their own `docker-compose.yml`, and proxy them all to the outside world via our `nginx-proxy` container.
- No other container can access containers within a default network created by docker-compose, only those inside the application's own `docker-compose.yml`. This makes life difficult for `nginx-proxy`.
- To work around this, we create a single network outside of Compose's infrastructure, and place our `nginx-proxy` container in that network. To do this, we need to define this network as `external` in the `docker-compose.yml` - this way Compose will not try to create the network itself, but just assign the containers it creates to the existing outside network.
- If we try to run this as is, Compose will error out telling us the external network doesn't exist.

We need to initially create the proxy network manually with:

```
sudo docker network create nginx-proxy
```

Now we can run our proxy container and SSL companion:

```
sudo docker-compose up -d
```

And continue with deploying at least one application behind the proxy to see that it actually works.

Running applications behind the proxy

For various application I run on the server, I create separate directories in `/home/olex`. Let's look at one of the simplest possible apps: a web server with some static content.

Single-container applications

To run such a server, one could use Apache, Lighttpd or another Nginx. I prefer to use things I know, and Nginx is plenty small and fast as a static content server, so that's what I go with. Here's the `docker-compose.yml` for such an "application", if you can even call it that:

```
version: '3'

services:
```



```
expose:
  - "80"
environment:
  - VIRTUAL_HOST=static.domain.tld
  - VIRTUAL_PORT=80
  - LETSENCRYPT_HOST=static.domain.tld
  - LETSENCRYPT_EMAIL=email@somewhere.tld
volumes:
  - /home/olex/static/files:/usr/share/nginx/html:ro
networks:
  - proxy
restart: unless-stopped

networks:
  proxy:
    external:
      name: nginx-proxy
```

Let's go through this point by point.

- There's a single service, `nginx`, running the tiny `nginx:alpine` image. We don't need more than that to serve some static content.
- There are no ports forwarded to the host, but port 80 is exposed - this is necessary so `nginx-proxy` can forward to it.
- Environment variables set all configuration values needed by `nginx-proxy` and `letsencrypt`:
 - `VIRTUAL_HOST` tells `nginx-proxy` under which domain should this container be reachable.
 - `VIRTUAL_PORT` tells `nginx-proxy` which port to forward requests to. In this case it's 80 as is usual for a HTTP server, it could be any other port - e.g. various Node.js apps typically serve up on port 3000.
 - `LETSENCRYPT_HOST` tells `letsencrypt` that this container's traffic should be SSL encrypted, and which domain to request a Let's Encrypt certificate for. It is a bit redundant to `VIRTUAL_HOST` (I can't think of a situation where you'd want them to be different), but that's how the proxy containers work.
 - `LETSENCRYPT_EMAIL` is used to tell Let's Encrypt where to send certificate expiration notices. Typically this never happens, because `letsencrypt` auto-renews certificates as necessary. If you get such an email and you haven't taken the service down, you'll at least know something's not working properly.



never be too careful).

- Container is assigned to the same external network as used above. This allows `nginx-proxy` to forward requests despite this being a separate Compose application, that would normally be completely isolated from other containers on the host.

Multi-container applications

Most applications you might want to deploy consist of several containers - say, a database, a backend service that does some processing, and a frontend. Of those containers, you may only want the frontend to be accessible from the outside, but the three should be able to talk to each other internally. A `docker-compose.yml` for such a setup might look something like:

```
version: '3'

services:

  frontend:
    image: frontend-image:latest
    expose:
      - "3000"
    environment:
      - VIRTUAL_HOST=awesomeapp.domain.tld
      - VIRTUAL_PORT=3000
      - LETSENCRYPT_HOST=awesomeapp.domain.tld
      - LETSENCRYPT_EMAIL=email@somewhere.tld
      - BACKEND_HOST=backend
    networks:
      - proxy
      - app

  backend:
    image: backend-image:latest
    environment:
      - DB_HOST=db
    networks:
      - app

  db:
    image: postgres:11
    volumes:
```



app

```
networks:
  proxy:
    external:
      name: nginx-proxy
  app:
```

Things of note here:

- Only the `frontend` container is put into the `proxy` external network, because it's the only one that needs to talk to `nginx-proxy`. It's also the only one with the environment variables for proxy configuration. The containers that are only used internally in the app don't need those
- `backend` and `db` containers are put in a separate `app` network, that is defined additionally under `networks`. This serves in place of a default network that Compose would normally create for all containers defined in the `docker-compose.yml` - except once you define any network manually, the default network doesn't get created anymore, and so our `backend` and `db` containers would be completely isolated. With the `app` network, they can talk to each other and to `frontend` normally.

As you can see, the setup is not very complicated. The only things to add to your normal `docker-compose.yml` of any multi-container app are the networks and environment variables for the frontend container.

Multiple proxied frontend containers

You can have multiple containers within a single Compose application that you want to make available via the proxy - for example, a public frontend and an administrative backoffice. You can configure as many containers inside one `docker-compose.yml` with the `proxy` network and necessary environment variables as you need. The only limitation is that every container needs to run on a separate `VIRTUAL_HOST`.

For example, you could run your public frontend at `awesomeapp.domain.tld` and the admin backend to that at `admin.awesomeapp.domain.tld`.

